

dstack

A Zero Trust Framework for
Confidential Containers

Phala Network
Verifiable & Confidential



Zero Trust Problem

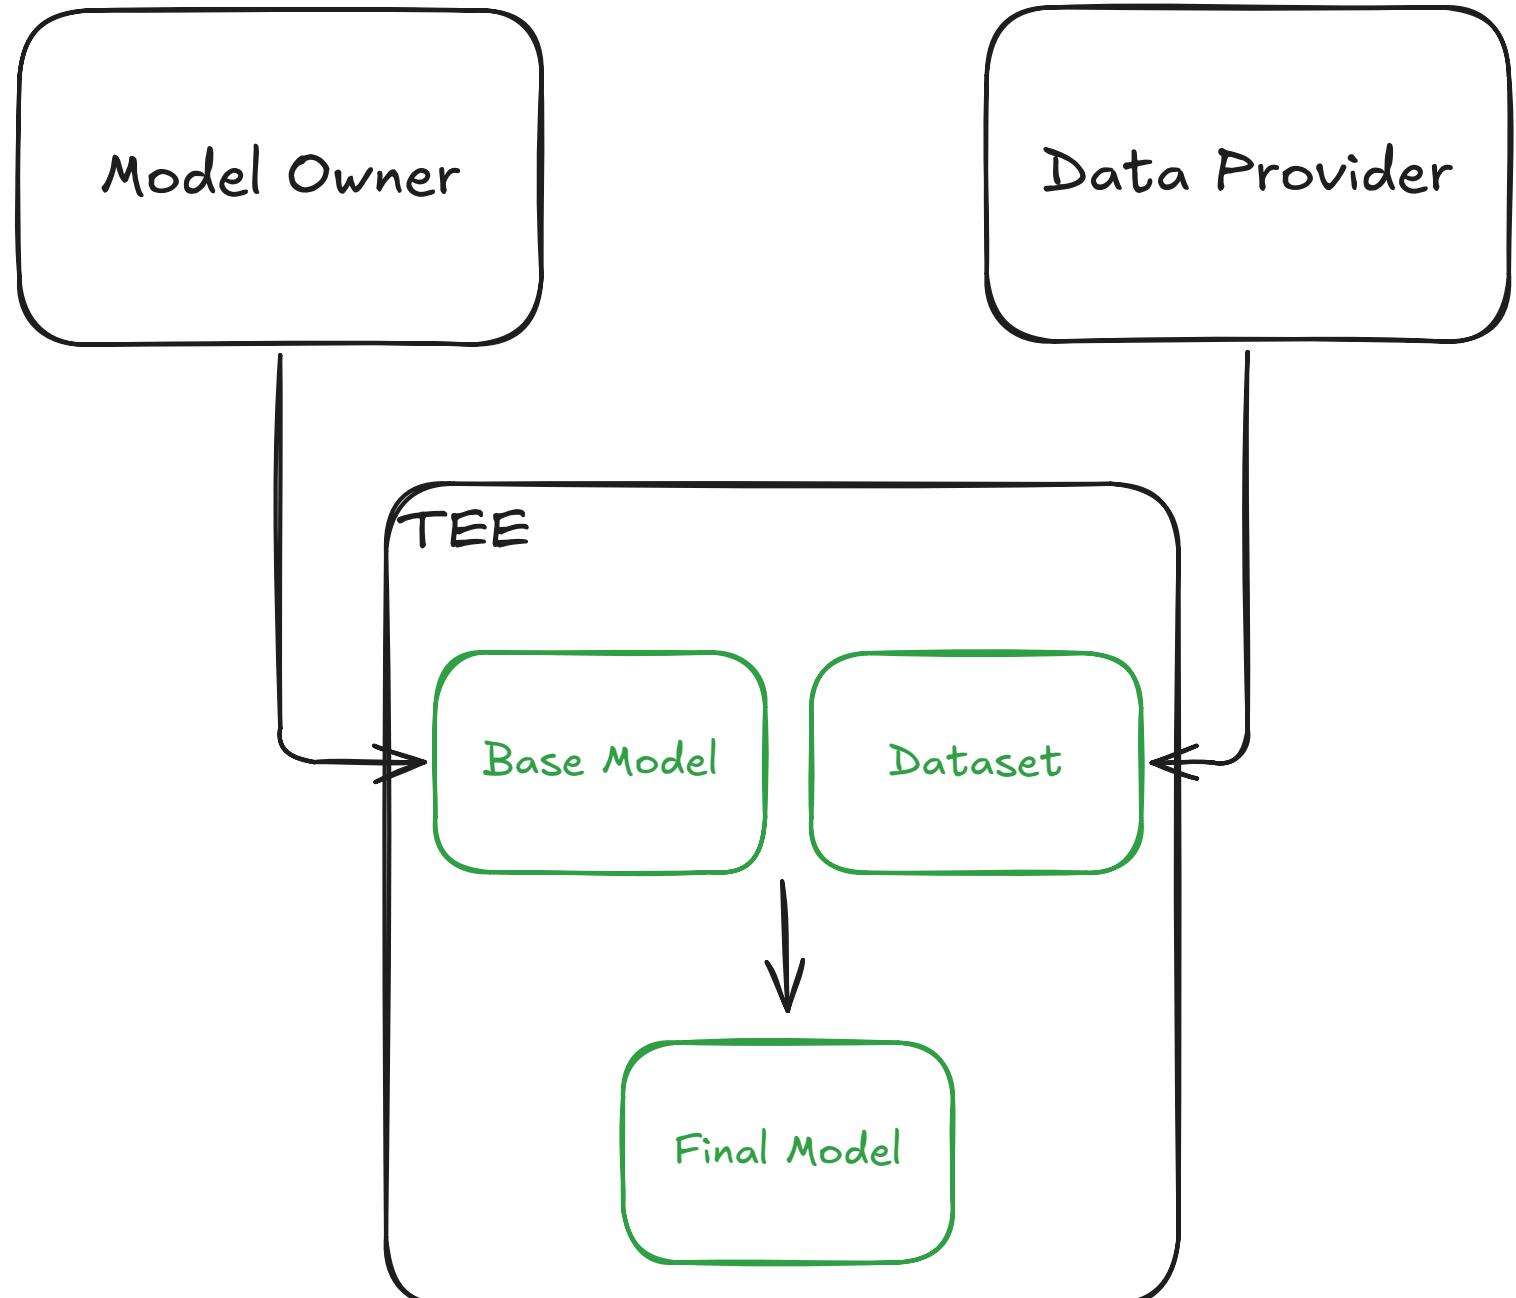
where nobody trusts each other

Requirements

Model parameters must remain confidential to Data Provider

Training data must remain confidential to Model Owner

Security guarantees must be verifiable by both parties



Zero Trust Problem

where nobody trusts each other

Threat Model

Model Owner is not trusted

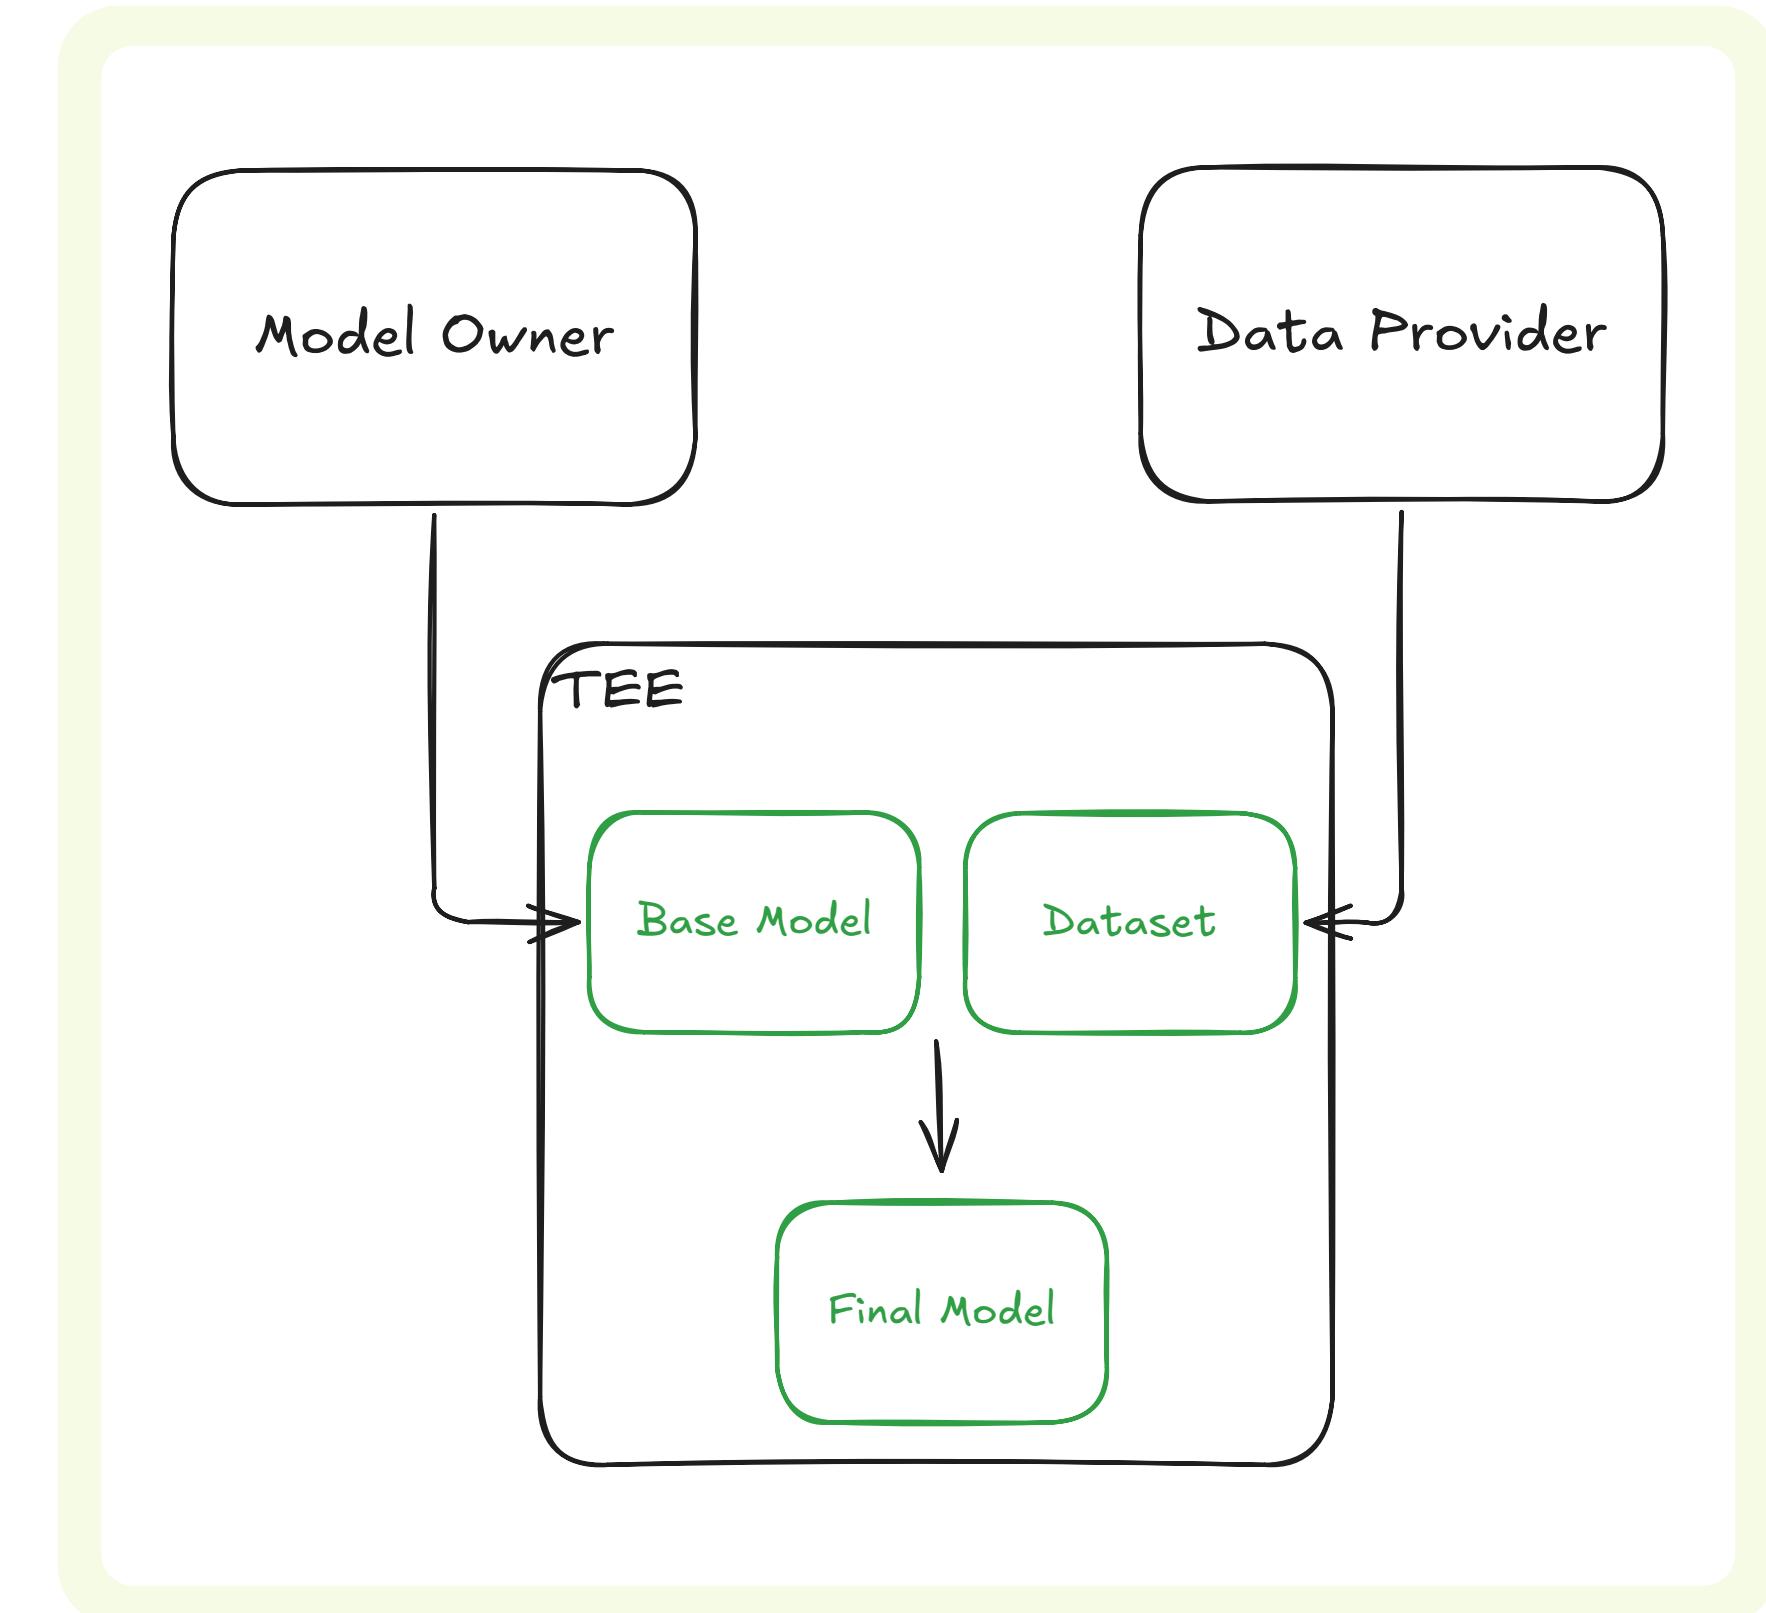
TEE vulnerabilities - hardware attacks, side channels, e.t.c.

Single point of failure - TEE can be down

What are missing in the current

TEE solutions

- ✗ Source Code Transparency
- ✗ Robust Key Management
- ✗ Continuous Security Assurance



TEE Vulnerabilities & Availability Gaps

Limitations of TEE

Known vulnerabilities compromise TEE confidentiality and integrity

Inconsistent key management across different TEE technologies

TEE shutdown = data loss: No cross-hardware recovery mechanism

TEE Solution	Key for Apps
Intel SGX	Hardware-bound keys
Intel TDX	No built-in API
AMD-SEV	No built-in API

TEE Vulnerabilities & Availability Gaps

The Recovery Problem

Core challenge: Restoring encrypted program state and data across different TEEs

Data backup is easy - **Key management is hard**

Need: Secure key recovery mechanism independent of specific hardware

✗ Extra Cost: Developers need to handle the encrypted data themselves, including **backup, key management and migration.**

TEE Solution	Key for Apps
Intel SGX	Hardware-bound keys
Intel TDX	No built-in API
AMD-SEV	No built-in API

Limited Verifiability

What can users learn from such a report?

- ✗ Is this really from the app I am interacting with?
 - ✗ Does it has backdoors in the code?
 - ✗ How will it process my data?

 **No Benefits:** TEE fails to bring **perceivable security and verifiability** to users.

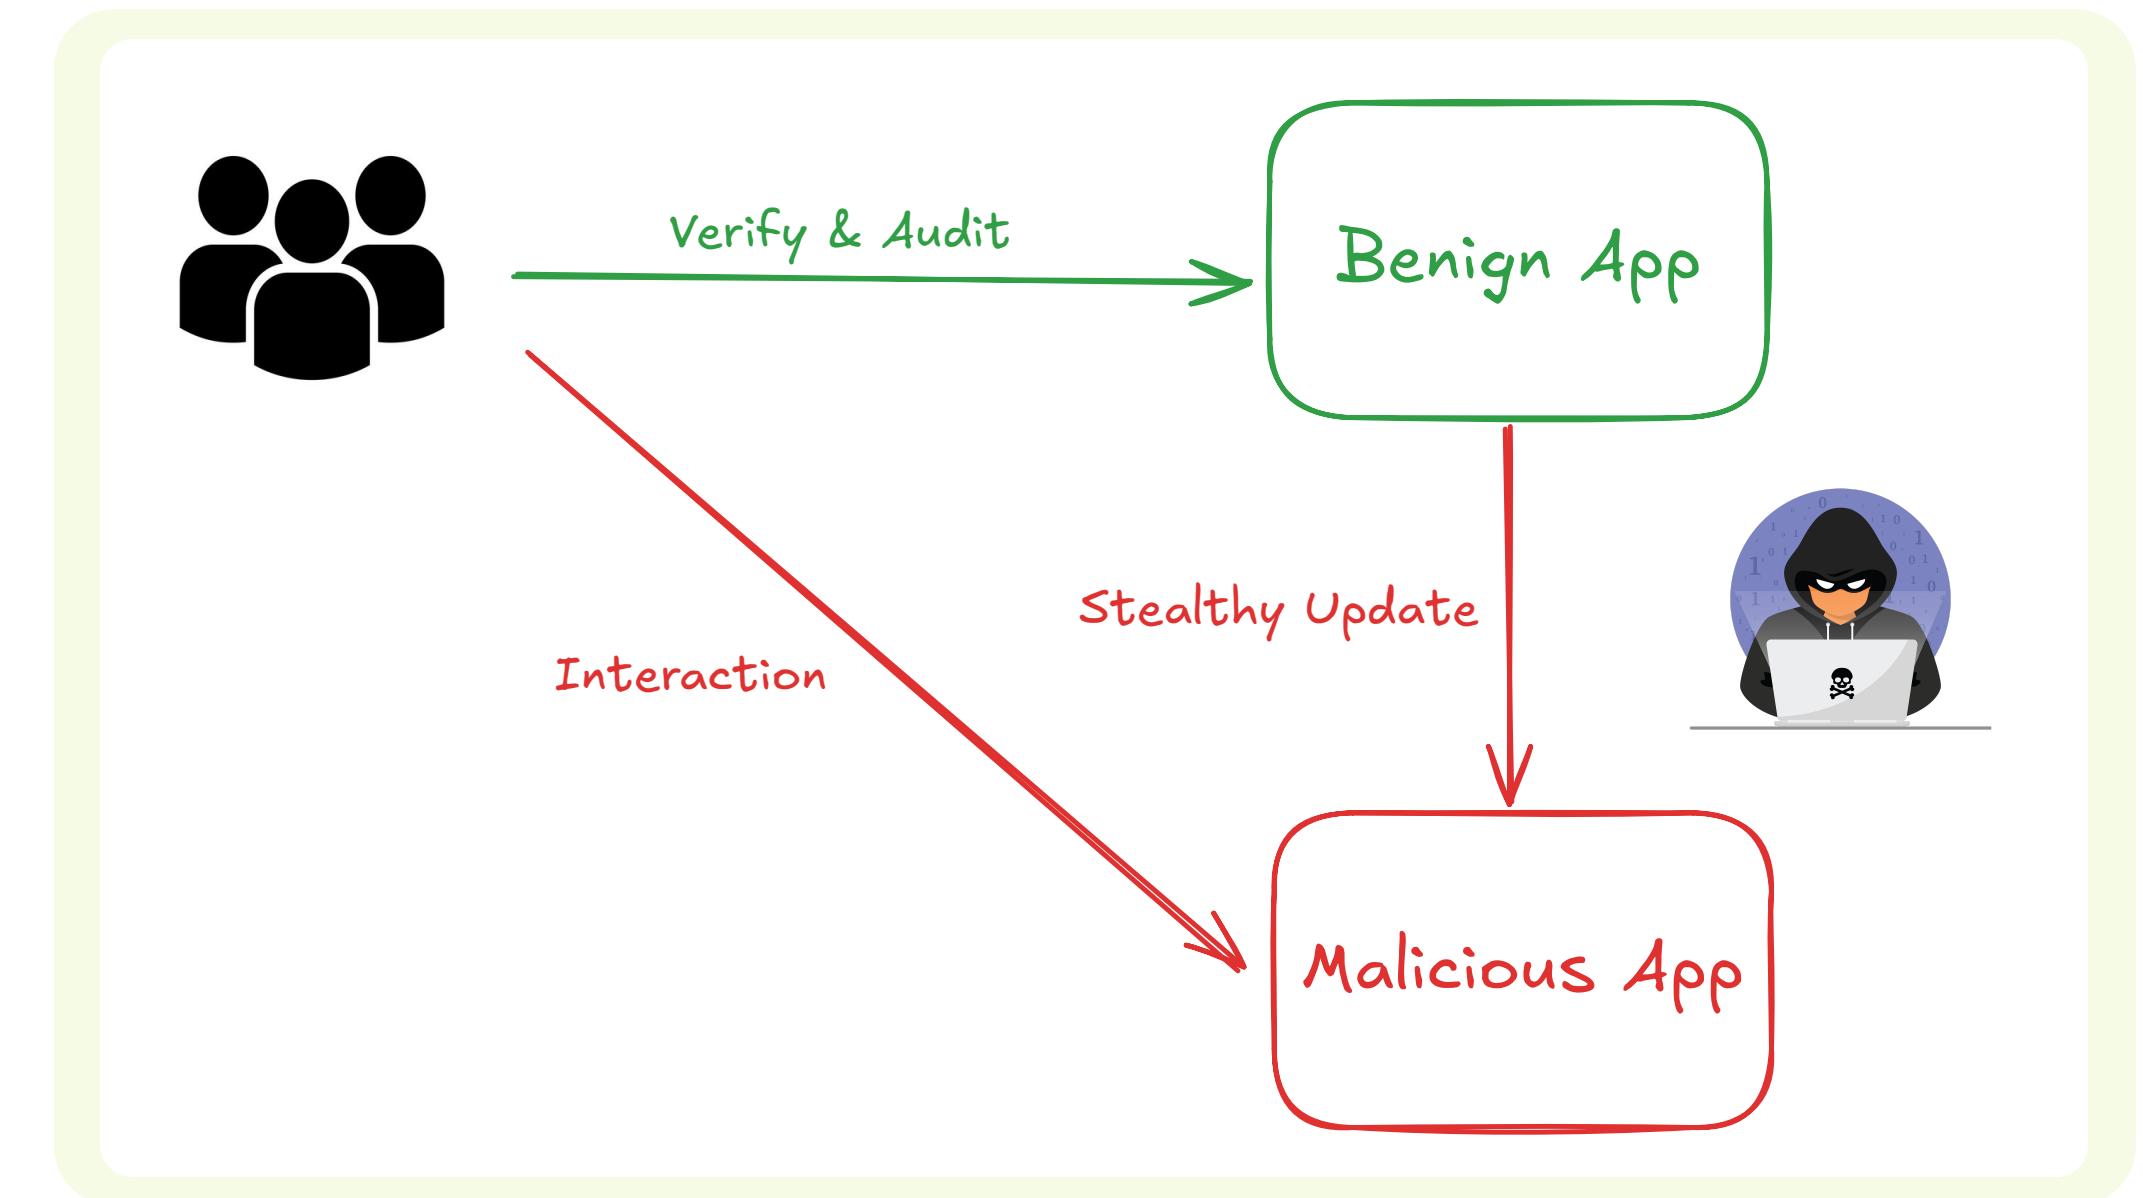
Beyond Current Trust Models: Addressing Malicious Administrators

Trust Model Evolved

There should not be "Administrators" any more

A **transparency log** of all application lifecycle changes is needed

✖ If the administrators can update the code **without being noticed**, the protection from TEE becomes meaningless.



Technical Challenges

Security Reliability: TEE vulnerabilities

Lack of Key Management creates
Single Point of Failure

Unrestricted Application Lifecycle Control

Attestation report with no context

Target

A Zero Trust Platform

Workaround for Attacks

Data & Service Availability

Code is Law

Fully Verifiable Chain-of-Trust

Zero Trust should come with Benefits and No Cost

What prevents developers from using TEE?

- Concerns about vulnerabilities and vendor lock-in
- Programming complexity: don't want to rewrite code
- Performance limitations and overhead
- TEE is invisible to customers



USE TEE



**CLAIM
TO USE TEE**

Dstack Innovations

Portable Confidential Container

Stateful CVM can be migrated between different TEE instances

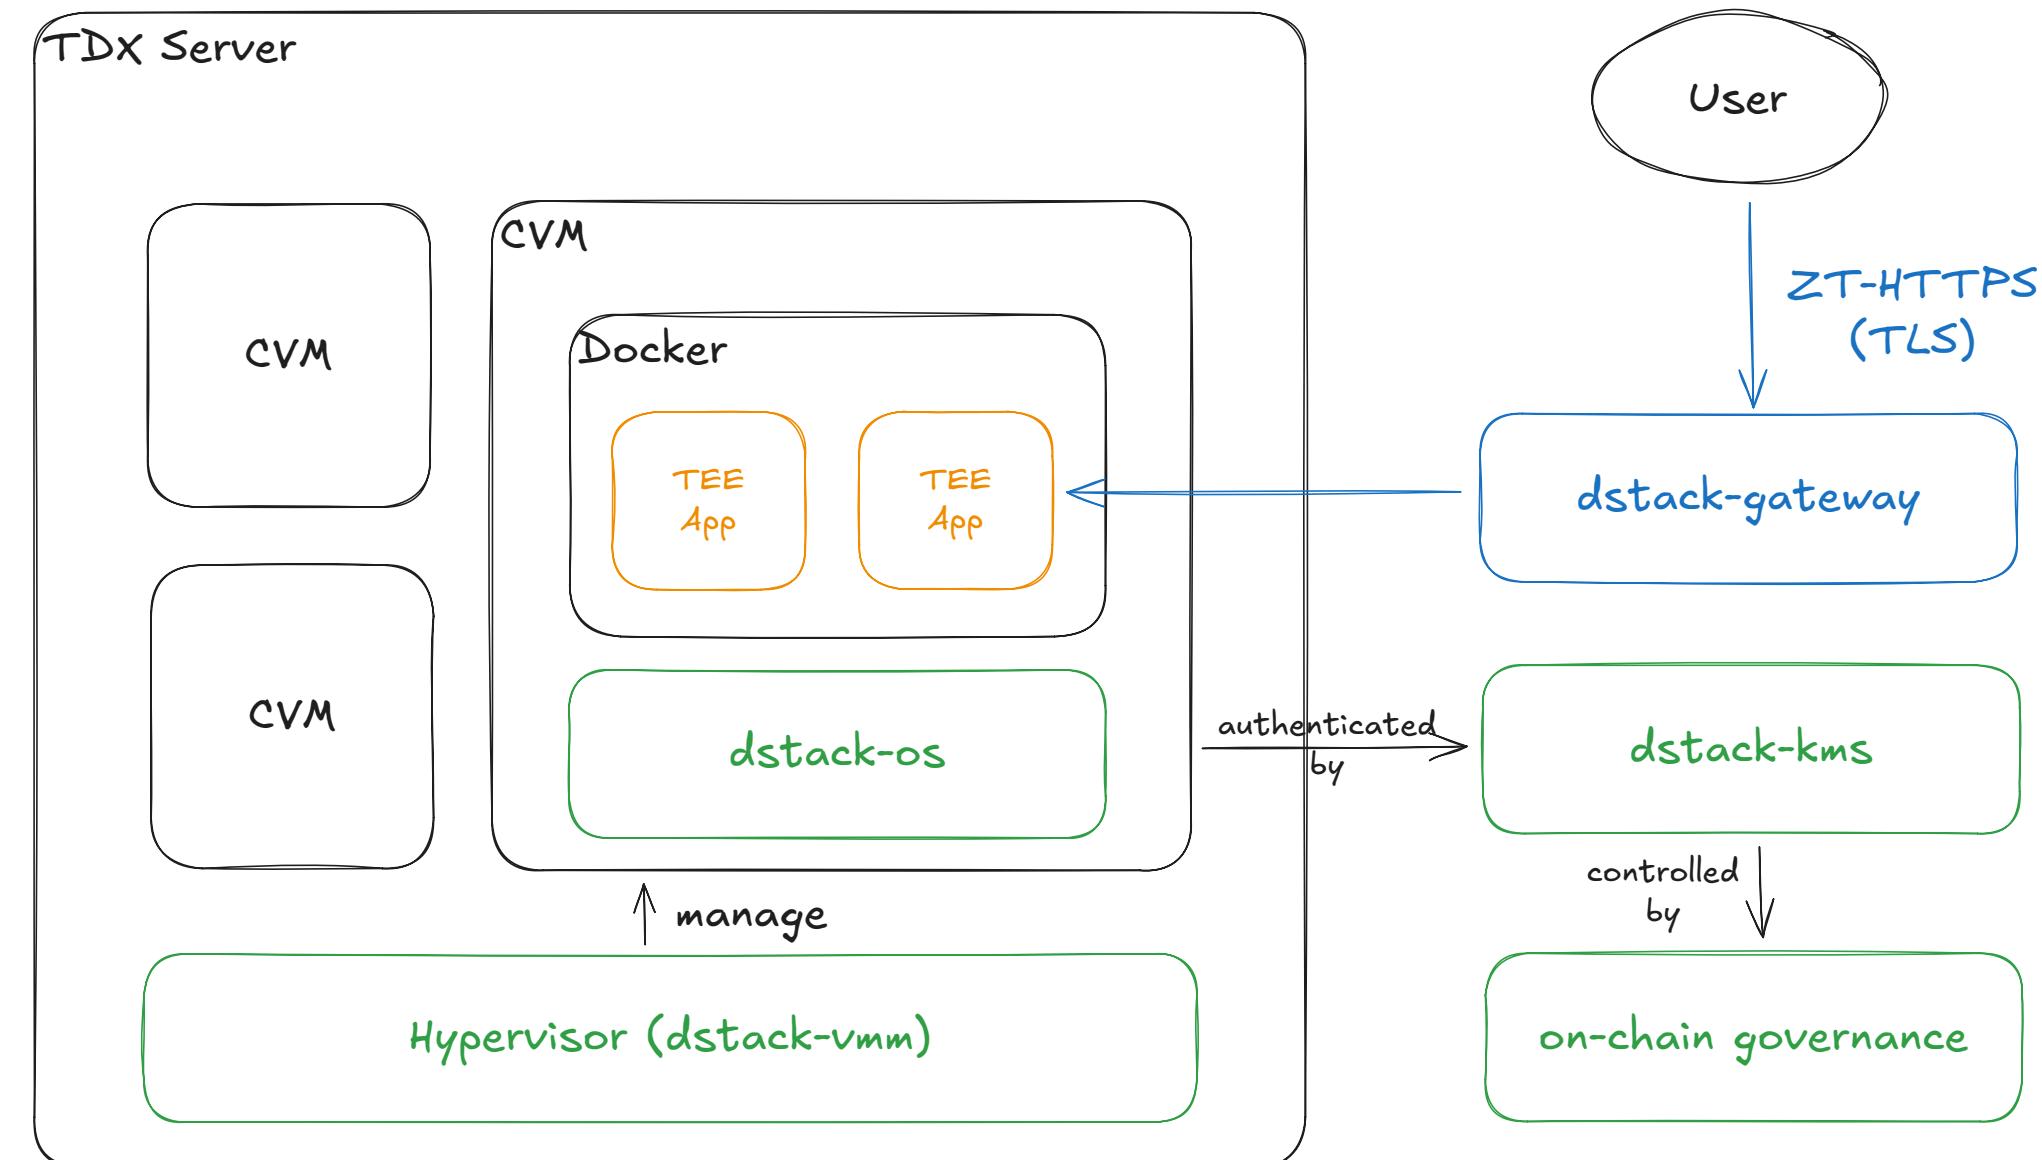
Code Management with Distributed Ledger

Transparent logs for all the application lifecycle events

Full Chain of Trust

Verifiable Domain Management with domain and certificate exclusively controlled by TEE

Source Code Verifiability for the whole system image



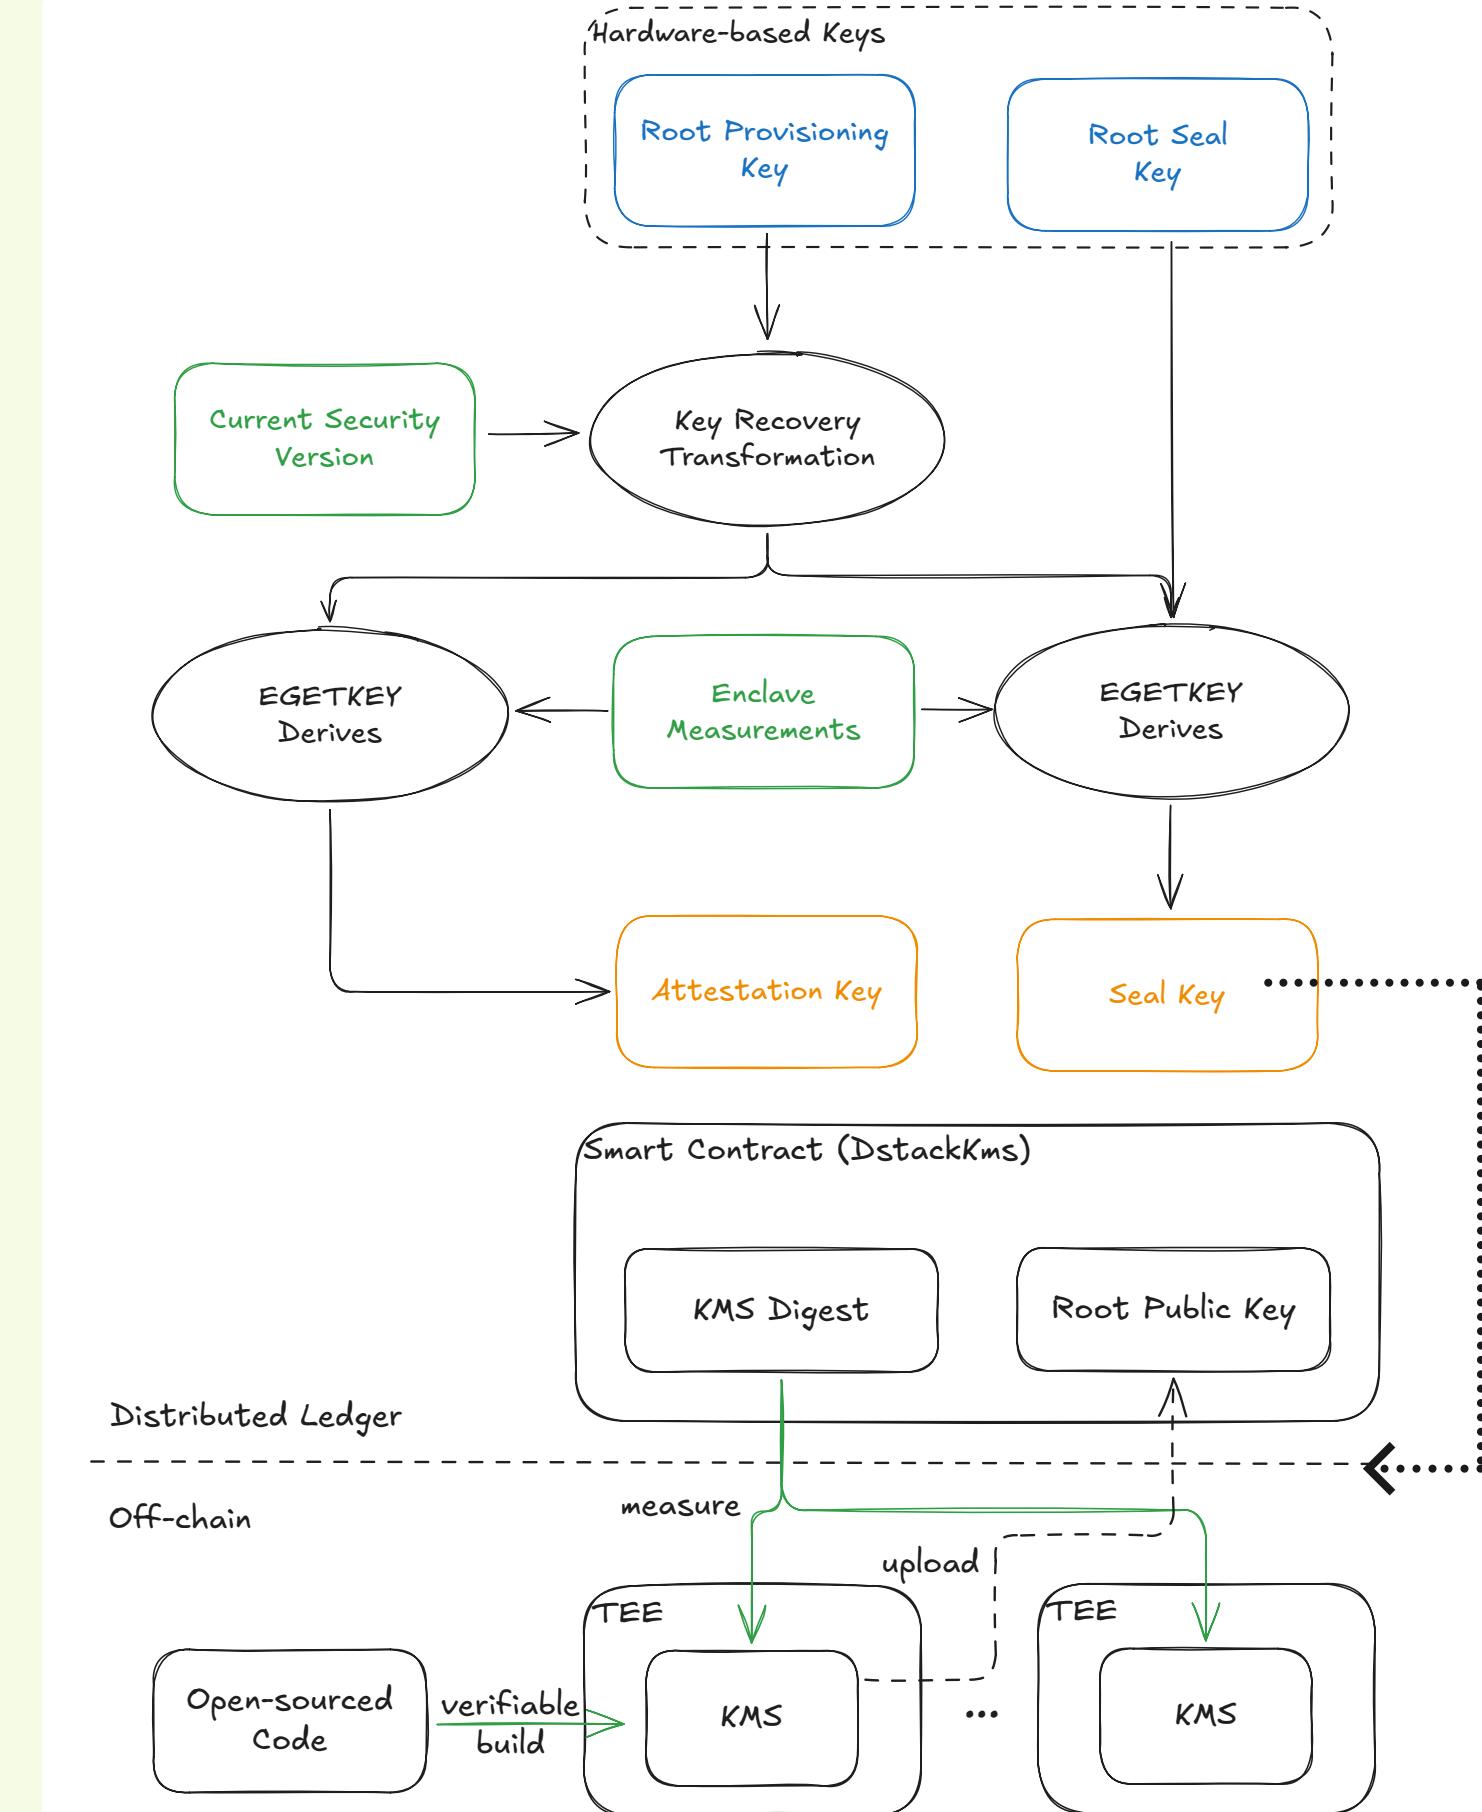
Dstack brings docker support to TEE CPU and GPU with features above

Portable Confidential Container

```
// The kms public RPC service.  
service KMS {  
    // Request the app key given the app id and tdx quote  
    rpc GetAppKey(GetAppKeyRequest) returns (AppKeyResponse);  
    // KMS key handover  
    rpc GetKmsKey(GetKmsKeyRequest) returns (KmsKeyResponse);  
    // Request the app environment encryption public key given the app id  
    rpc GetAppEnvEncryptPubKey(AppId) returns (PublicKeyResponse);  
    // Request the KMS instance metadata for use as a probe and health check.  
    rpc GetMeta(google.protobuf.Empty) returns (GetMetaResponse);  
    // Request the temporary CA certificate and key  
    rpc GetTempCaCert(google.protobuf.Empty) returns (GetTempCaCertResponse);  
    // Sign a certificate  
    rpc SignCert(SignCertRequest) returns (SignCertResponse);  
    // Clear the image cache  
    rpc ClearImageCache(ClearImageCacheRequest) returns (google.protobuf.Empty);  
}
```

Generate Secret Keys with Independent Key Management Service

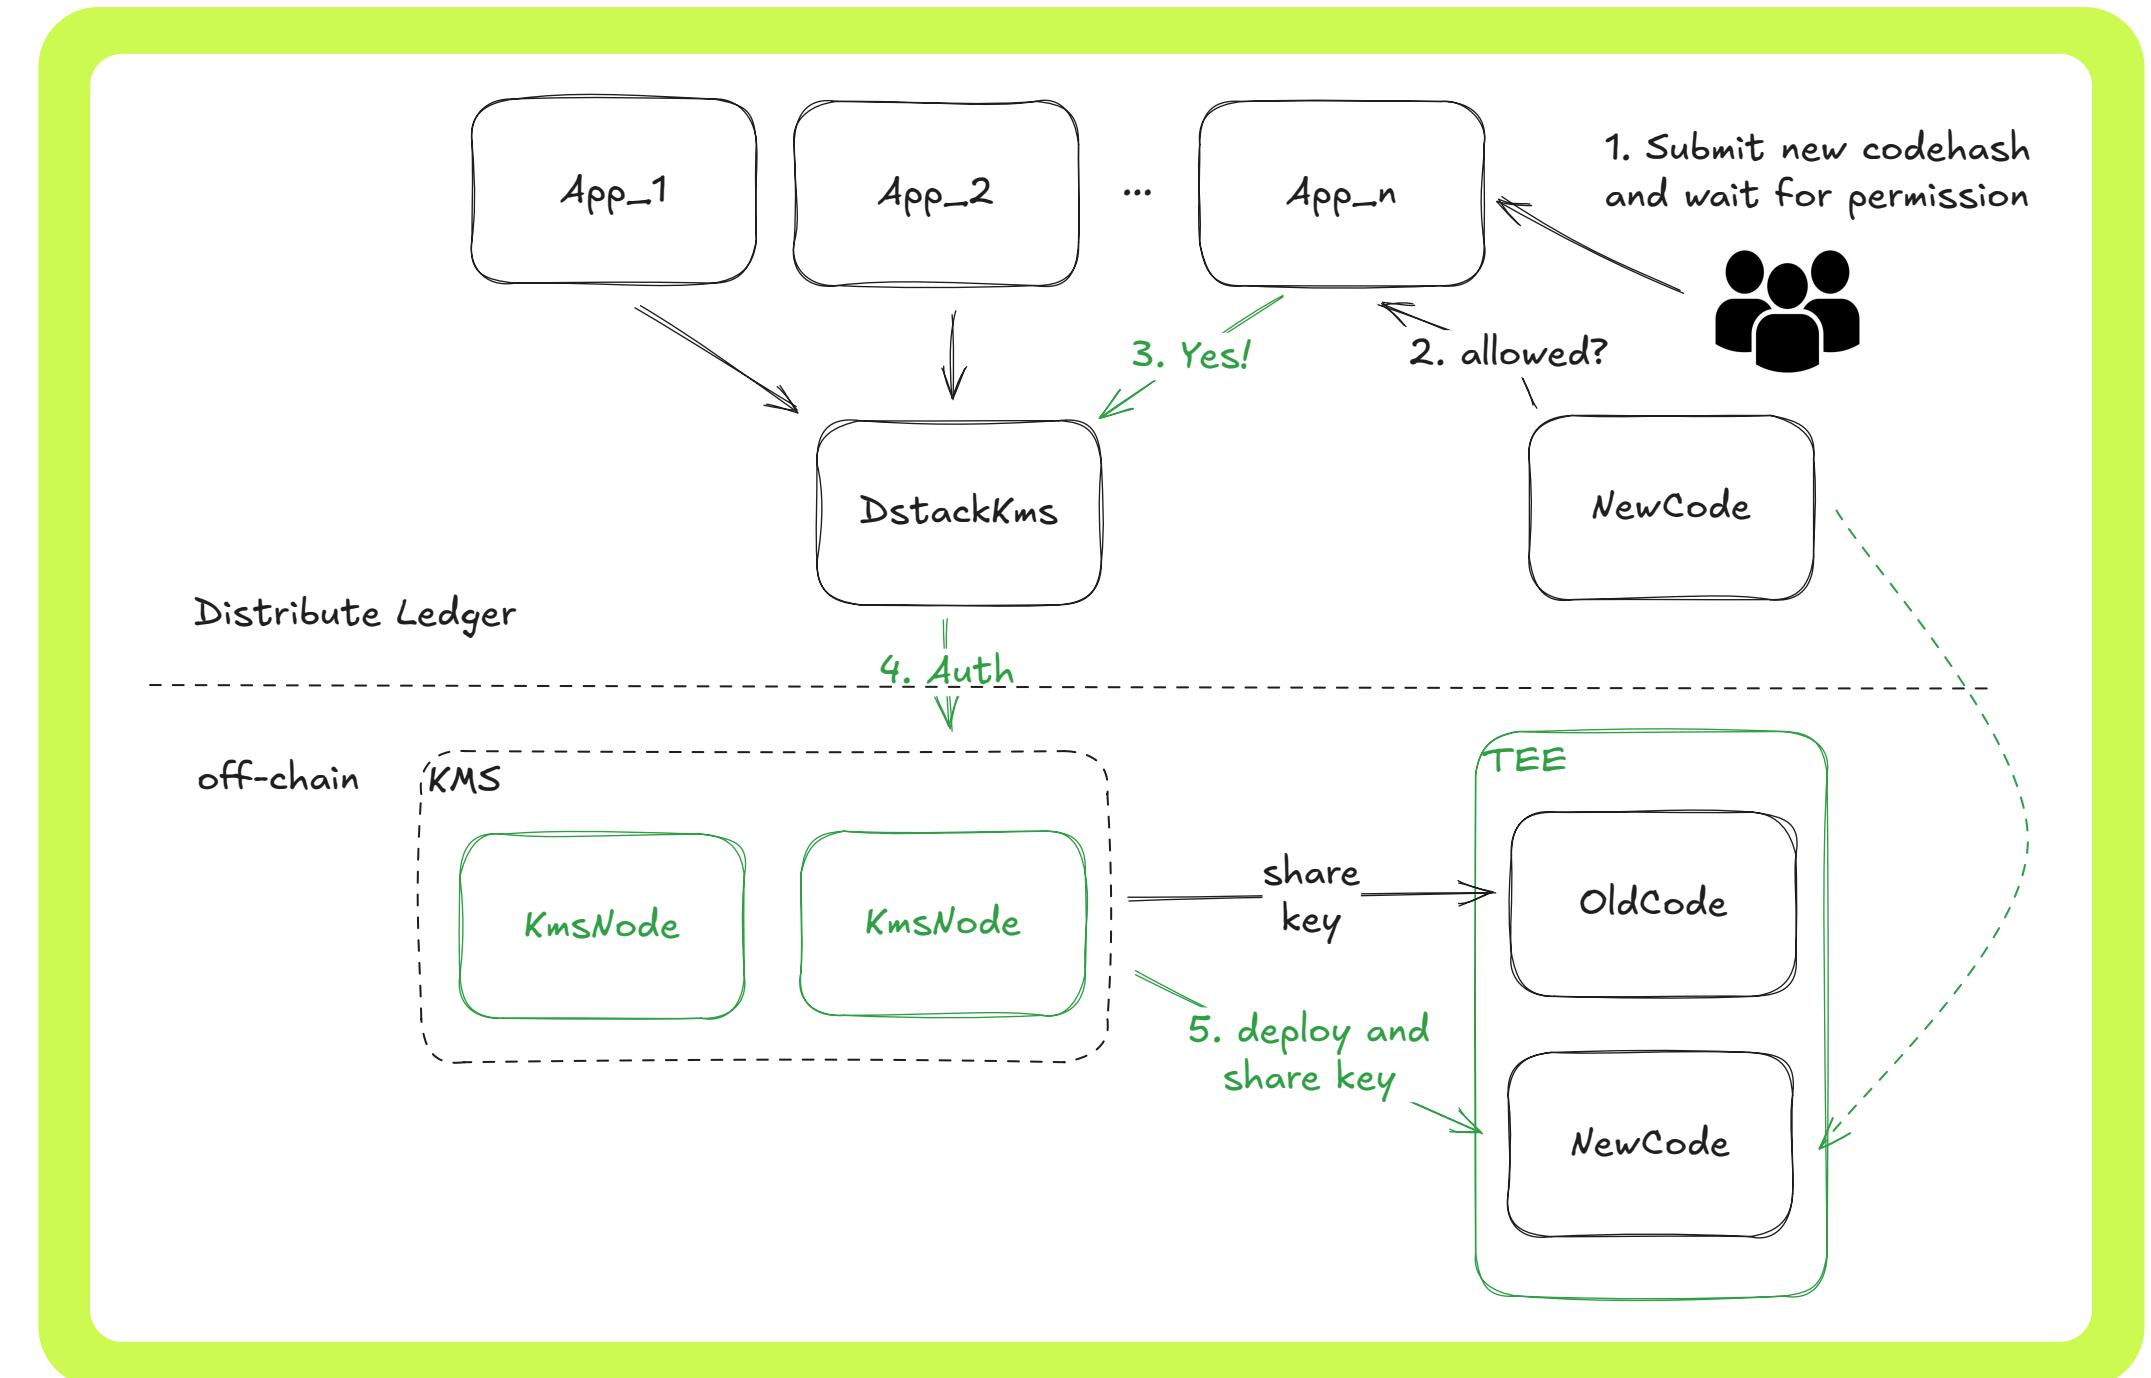
The KMS becomes the new Root-of-Trust of the whole system
By default, KMS is powered by a P2P TEE network



Decentralized Code Management

Customizable management
enforced through KMS

Without KMS approval, the TEE hardware
will not be able to decrypt all its data



Verifiable Domain Management

with TEE exclusively controls a domain

TEE certificate generation

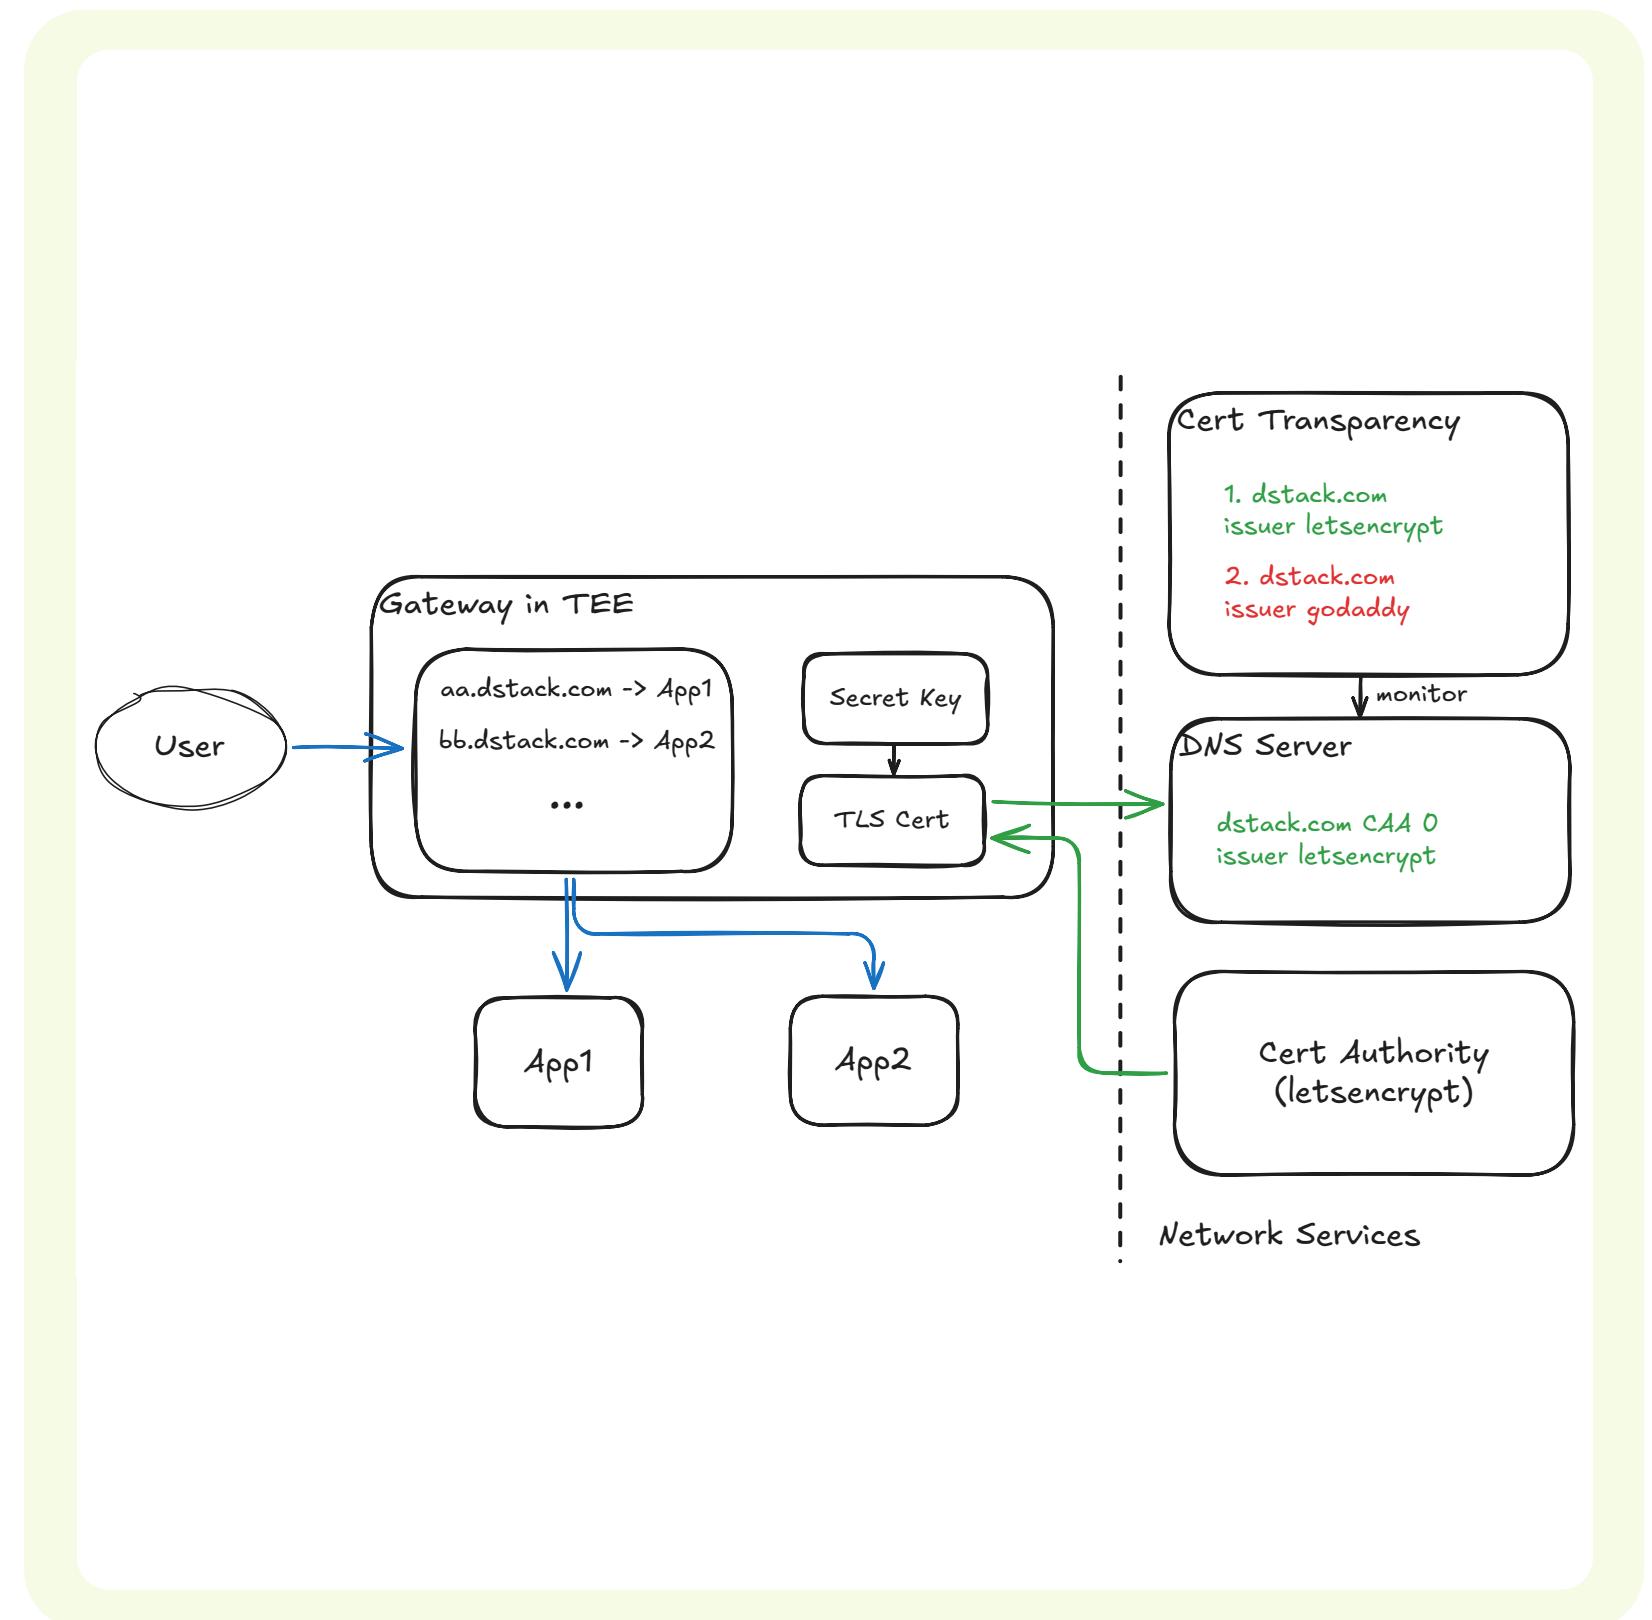
Attestation report will be generated to show an account_uri is bound to **TEE-controlled key**

Ownership monitoring with CAA records and CT

Certificate Authority Authorization restricts
which certificate authorities can issue certificates for
managed domains

Certificate Transparency logs provide an **immutable audit trail of certificate issuance history**

crt.sh ID	Logged At	Not Before	Not After	Common Name	Matching Identities	Issuer Name
18790913612	2025-06-03	2025-06-03	2025-09-01	*.dstack-testnet.phala.network	*.dstack-testnet.phala.network	C=US,O=Let's Encrypt,CN=E6
18790970093	2025-06-03	2025-06-03	2025-09-01	*.dstack-testnet.phala.network	*.dstack-testnet.phala.network	C=US,O=Let's Encrypt,CN=E6
18719669040	2025-05-30	2025-05-30	2025-08-28	*.dstack-testnet.phala.network	*.dstack-testnet.phala.network	C=US,O=Let's Encrypt,CN=E6
18719668685	2025-05-30	2025-05-30	2025-08-28	*.dstack-testnet.phala.network	*.dstack-testnet.phala.network	C=US,O=Let's Encrypt,CN=E6



Walkthrough of Chain-of-Trust

Audit of dstack Components

Scope

Low-level libraries and tooling, including remote attestation components, guest agent service, and utilities
Image-related files for **dstack-OS**

Findings

A **high-severity issue** with the OVMF build configuration, now fixed with a more secure implementation

Medium-severity issues related to terminal binaries, symbolic link handling, and environment variable protection

Several **low-severity findings** and documentation improvement recommendations



Audit of dstack

Date: May 26th, 2025

ID	COMPONENT	NAME	RISK
#00	meta-dstack/ovmf	VMM is Currently Trusted in OVMF Build	High
#01	meta-dstack recipes	Terminal Binaries Present in Production Dstack Image	Medium
#02	dstack-util system setup	Host Can Pass Symbolic Links To Shared Folder With Guest	Medium
#03	dstack-util	Env Injection via Unauthenticated Shared Files	Medium
#04	app-compose service	Pre-Launcher Code Can Be Used To Leak Secrets on Default KMS	Medium
#05	meta-dstack	qemu-guest-agent is Present in Production	Medium
#06	dcap-qvl	Incomplete TD Under Debug Checks	Medium
#07	app-compose service	Unchecked Container Image Digest	Medium
#08	guest-agent	Unrestricted Exposure of stdout/stderr From CVM Docker Containers	Low
#09	dstack-util	Incomplete Measurement of CVM Configuration Files	Low
#0a	*	Underdocumented Root of Trust and Vendoried Attestation Code	Low
#0b	dcap-qvl	Lack of Revocation Checks in Quote Verification Library	Low
#0c	meta-dstack	Lack of Documentation on Design and Hardening Decisions in meta-dstack Layer	Informational

First Benchmark for Running LLM on TEE GPU

Conclusions

Minimal computational overhead within the GPU, the overall performance penalty is primarily attributable to **data transfer**.

For most of the LLM queries, the overhead remains **below 7%**.

Larger models and longer sequences experiencing **nearly zero overhead**.

Confidential Computing on NVIDIA Hopper GPUs: A Performance Benchmark Study

Jianwei Zhu, Hang Yin, Peng Deng[†], Aline Almeida[‡], Shunfan Zhou
Phala Network, [†]Fudan University, [‡]io.net
{jianweiz, hangyin, shelvenzhou}@phala.network,
[†]pdeng21@m.fudan.edu.cn, [‡]aline@io.net

November 6, 2024

Abstract

This report evaluates the performance impact of enabling Trusted Execution Environments (TEE) on NVIDIA Hopper GPUs for large language model (LLM) inference tasks. We benchmark the overhead introduced by TEE mode across various LLMs and token lengths, with a particular focus on the bottleneck caused by CPU-GPU data transfers via PCIe. Our results indicate that while there is minimal computational overhead within the GPU, the overall performance penalty is primarily attributable to data transfer. For the majority of typical LLM queries, the overhead remains below 7%, with larger models and longer sequences experiencing nearly zero overhead.

Acknowledgments

We would like to express our gratitude to the io.net [io.] and IOG Foundation [Fou] for their generous grant, which made this research possible. We also extend our thanks to Engage Stack [Sta], the cloud service provider, for providing the necessary hardware and technical support.

1 Introduction

Trusted Execution Environments (TEEs) are increasingly important in machine learning and AI due to growing security requirements in both enterprise and decentralized applications [SAB15, MSM⁺18, AKKH18]. The introduction of TEE-enabled GPUs, such as the NVIDIA H100 and H200, adds an extra layer of protection for sensitive data but may impact performance. Understanding these trade-offs, particularly for large-scale machine learning tasks, is crucial for adopting TEE in high-performance AI applications [YMY⁺22, WO24].

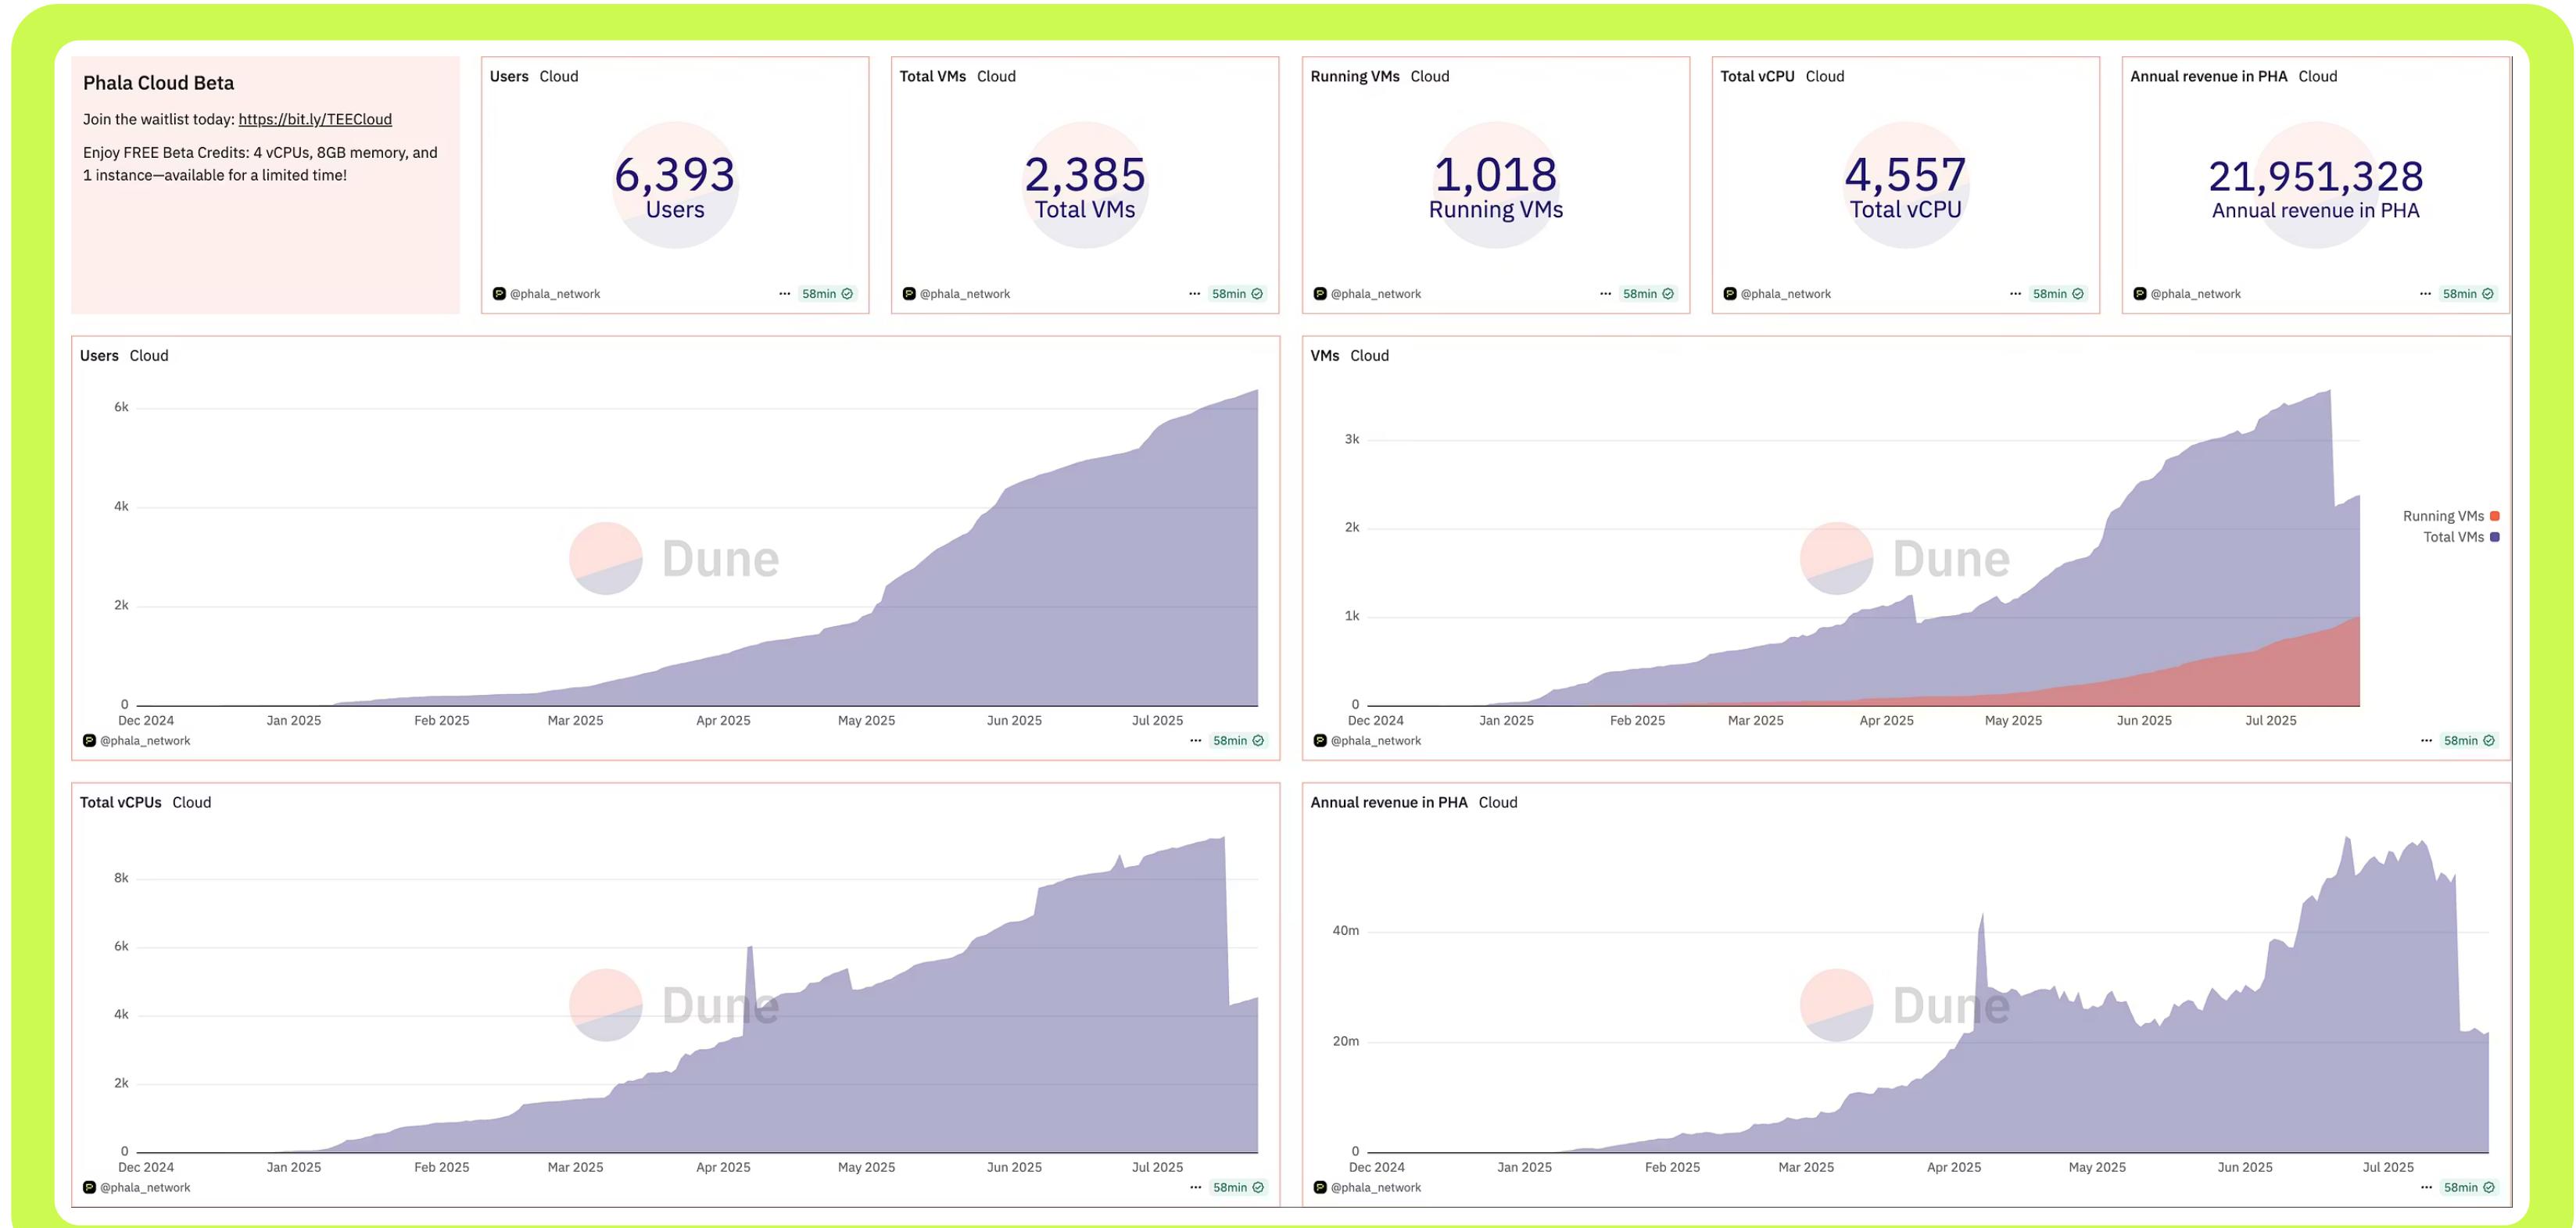
This report quantifies the performance overhead of enabling TEE on the NVIDIA Hopper architecture GPUs during LLM inference tasks, identifying where the overhead arises and under what conditions it can be minimized.

2 Background

2.1 Trusted Execution Environment

A TEE is a hardware-based security feature that isolates computations, preventing unauthorized access and tampering, even from the operating system or the physical hardware owner. As the core technology enabling Confidential Computing, TEEs create secure enclaves where sensitive data and code are processed with encryption, ensuring confidentiality and integrity even if the broader system is compromised [SAB15]. Traditionally implemented in CPUs, TEE technology was extended to GPUs by NVIDIA in 2023, enabling tamper-proof and confidentiality-preserving computation inside the GPU with minimal performance penalty [DGK⁺23].

Usage Statistics



Project Status

Contributors



Phala Network

Core Contributor



Near AI

Private ML SDK



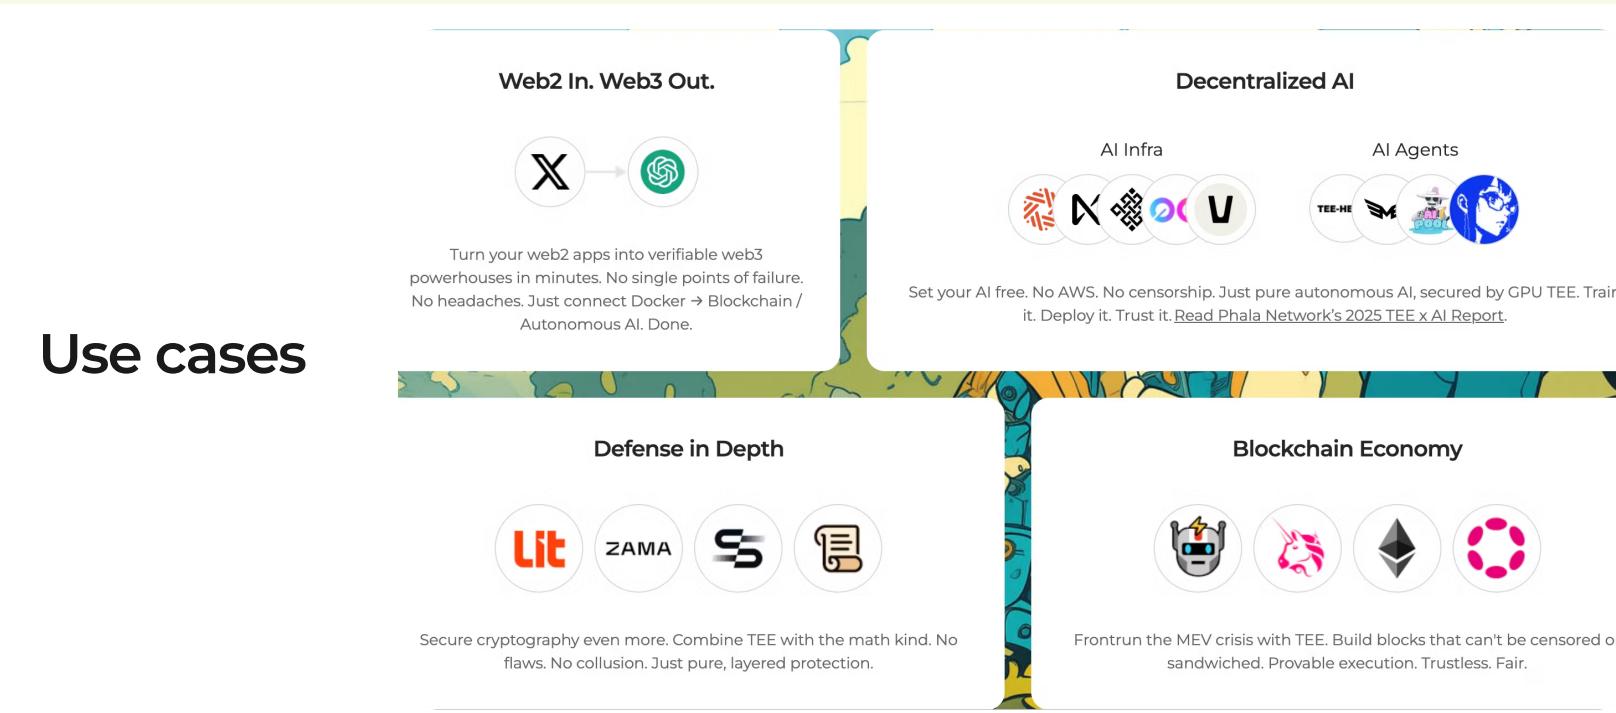
Prof. Andrew Miller

Core Contributor



Flashbots

Contributor



dstack

Deploy containerized apps to TEE with end-to-end security in minutes.

stars 223 license Apache-2.0 Community Ask DeepWiki

59 Contributions in the Last 30 Days

0 Opened/Closed Issue Ratio

-1 (-100%) past month

16 Pull Requests Opened

-14 (-46.67%) past month

73 Commits

-36 (-33.03%) past month

Issues

Opened
 Closed

Pull Requests

Opened
 Closed

Pushes & Commits

Pushes
 Commits

A [Live Community dstack Applications](#) explorer is available

Future Plan



Distributed System

K8s support

High availability

E2EE network



Compatibility

AMD SEV-SNP support

Nitro support



Security

Safer KMS with MPC + TEE

Attack detection with honeypot

Sandboxing

A [comprehensive version](#) is available



PHALA

Thank you !

Dr. Shunfan (Shelven) Zhou

Lead Researcher

shelvenzhou@phala.network

<https://cloud.phala.network/>

Comparisons between dstack and Confidential Containers

Similar Security Design Principle

Independent KMS to manage the secrets for running TEE workers / pods

Different Requirements on Verifiability

Dstack is designed with verifiability as a core requirement

Different Approaches to Usability

Dstack prioritizes **ease of use** with built-in security features

CoCo adopts a **minimalist approach**, leaving these to application developers

	dstack	Coco
Reproducible Build	All components: os images, kms and gateway	No
Code Audit	Core components	No
Measurement Target	Docker Image / Source Code	Sandbox and Security Policy
Attestation Generation	On request	Static during Sandbox Creation
Encrypted Network Connection	Built-in Support	Instruction only
Encrypted Disk Support	Built-in Support	Fs API available
Trust Center	Yes	No
Lifecycle Transparency Log	Built-in Support	No